



Prototype networks for few-shot learning

Fabian Greavu, Machine Learning exam @ unifi



Summary

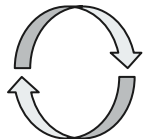
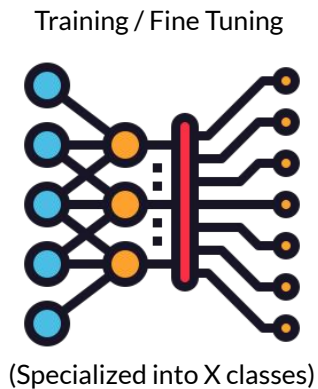
- ❖ Introduction
 - Challenge
 - Few-shot learning
- ❖ Project workflow
 - Datasets
 - Prototypical Networks
 - Centroids
 - NC, NS, NQ
 - Loss, Optimizer, scheduler
 - Training skeleton
 - Code
- ❖ Experiments
 - Omniglot
 - Mini Imagenet
 - Flowers102
 - Comparison
- ❖ Conclusion



Intro

Challenge

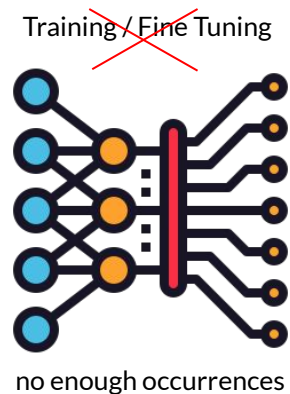

'classic' Classification Task



Deploy



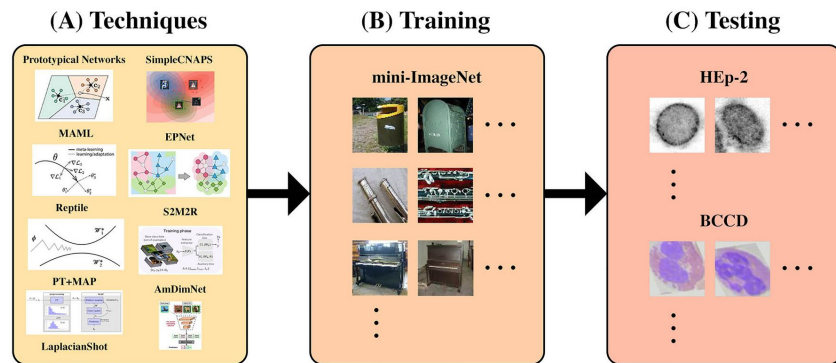
Our Challenge Classification Task



...

Few-shot learning


- A **meta-learning** technique to 'learn to learn'
- Uses **small amount** of occurrences
- Can perform with **high accuracies**




Dataset



Learn to learn



Learn + Deploy



Meta-learning definition

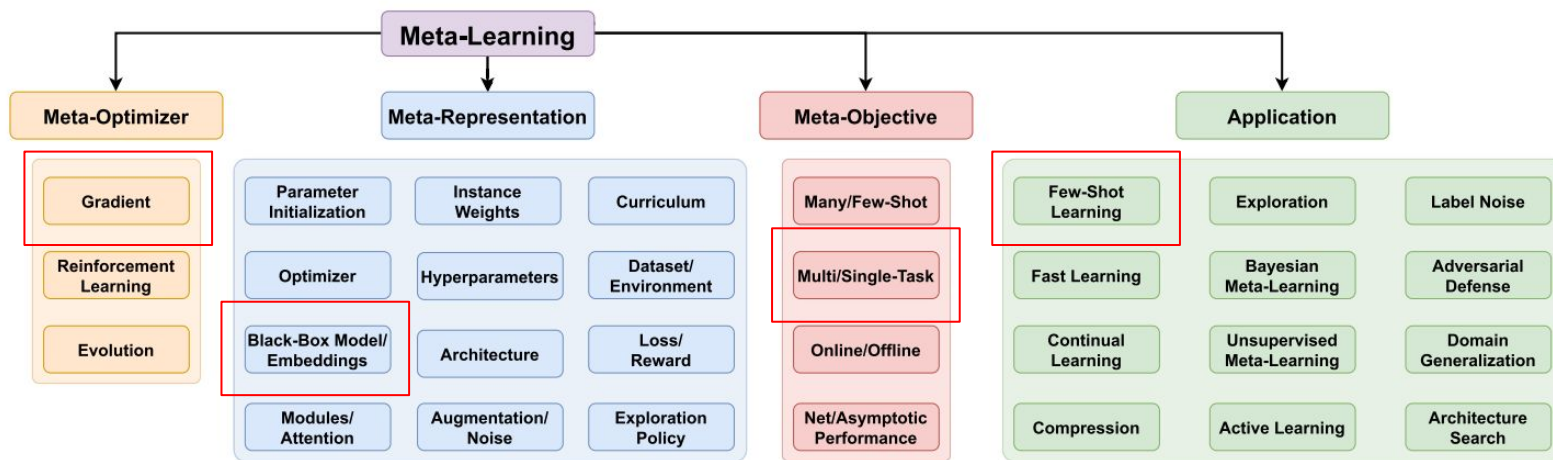


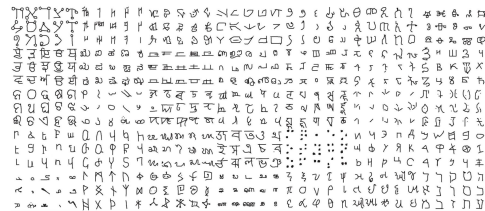
Fig. 1. Overview of the meta-learning landscape including algorithm design (meta-optimizer, meta-representation, meta-objective), and applications.



Project workflow

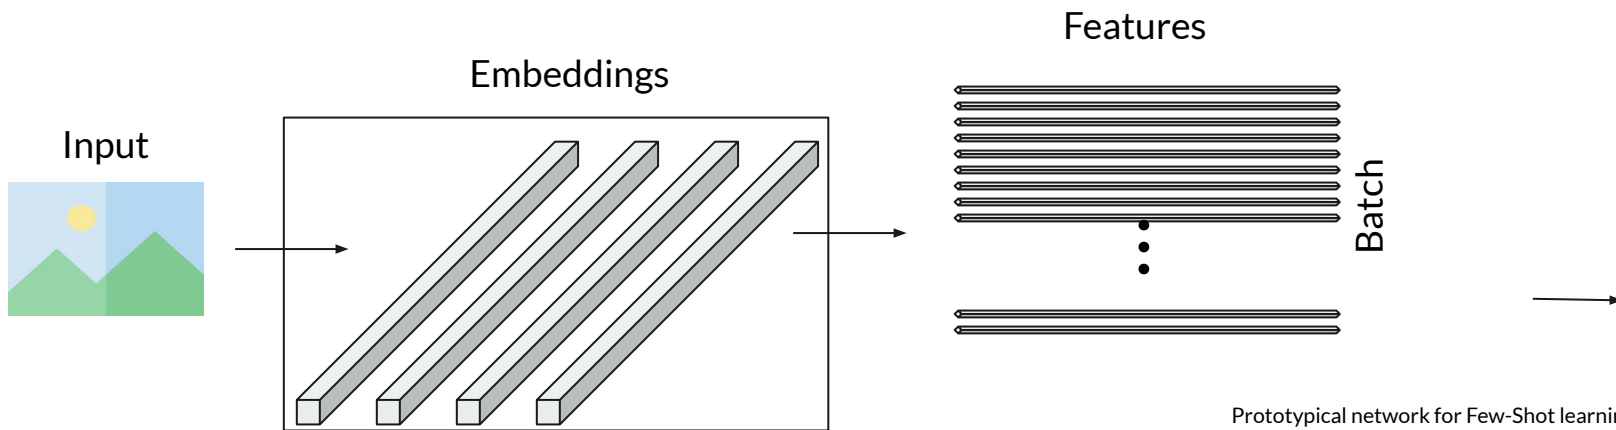
Datasets

- Omniglot: 1623 handwritten characters
 - 80 images per class
 - classes: (1032 train, 172 val, 464 test)
- Mini Imagenet: 100 objects
 - 600 images per class
 - classes: (64 train, 16 val, 20 test)
- Flowers102: 102 flowers
 - 40-120 images per class
 - classes: (64 train, 16 val, 22 test)



Prototypical Networks

- Extracts features with a neural network
 - 4 CNN blocks: Conv out=64, ks=3 / BatchNorm2D / ReLu / MaxPool2D
- Learning: embeddings learning
 - Trained 100 episodes/iterations per epochs (200 ep) to learn embeddings



Prototypical Networks - centroids

- Use a small part of output as support, other as query
- Centroids are mean(support)
- Calculate distances between query and centroids
- Calculate loss as mean of log_softmax of negative distances

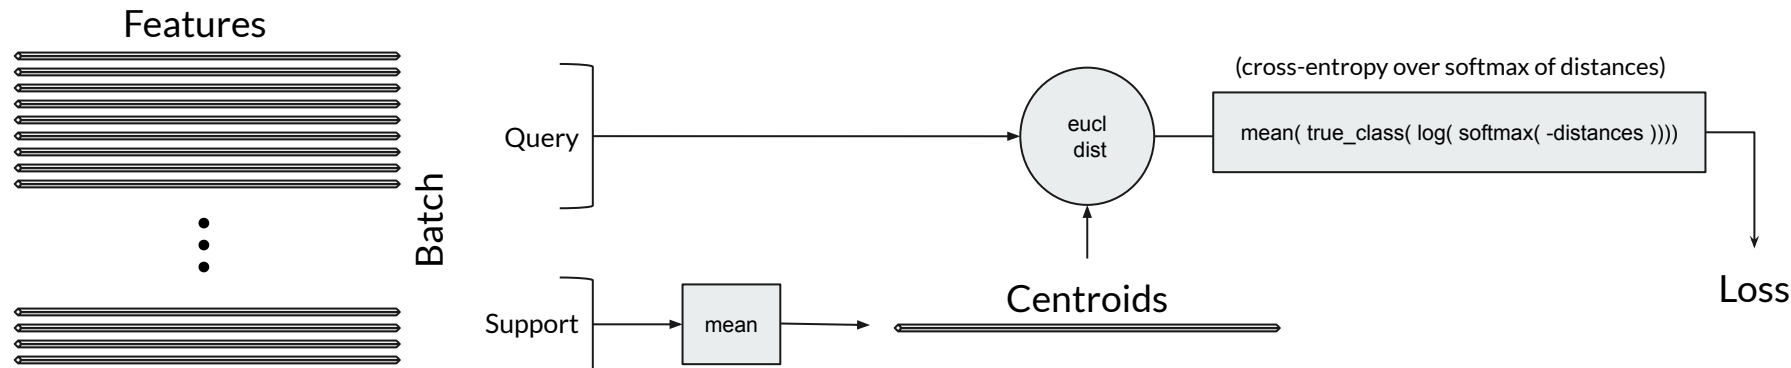
Prototypical networks compute an M -dimensional representation $\mathbf{c}_k \in \mathbb{R}^M$, or *prototype*, of each class through an embedding function $f_\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ with learnable parameters ϕ . Each prototype is the mean vector of the embedded support points belonging to its class:

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}_i, y_i) \in S_k} f_\phi(\mathbf{x}_i) \quad (1)$$

Given a distance function $d : \mathbb{R}^M \times \mathbb{R}^M \rightarrow [0, +\infty)$, prototypical networks produce a distribution over classes for a query point \mathbf{x} based on a softmax over distances to the prototypes in the embedding space:

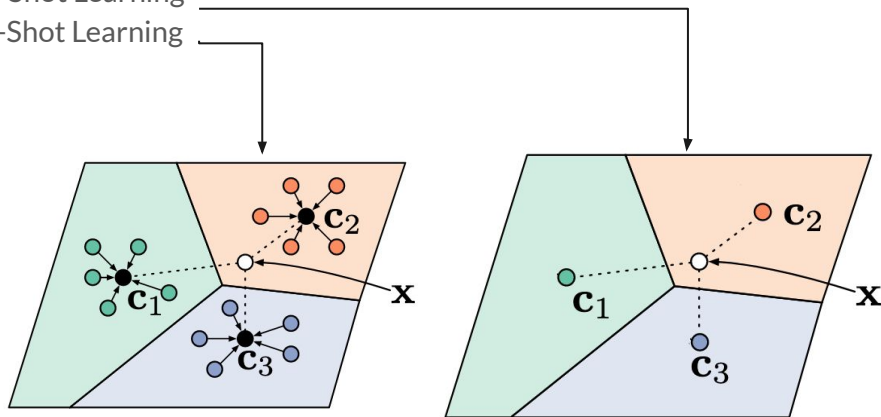
$$p_\phi(y = k | \mathbf{x}) = \frac{\exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_k))}{\sum_{k'} \exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_{k'}))} \quad (2)$$

Learning proceeds by minimizing the negative log-probability $J(\phi) = -\log p_\phi(y = k | \mathbf{x})$ of the true class k via SGD. Training episodes are formed by randomly selecting a subset of classes from the training set, then choosing a subset of examples within each class to act as the support set and a subset of the remainder to serve as query points. Pseudocode to compute the loss $J(\phi)$ for a training episode is provided in Algorithm 1.



Prototypical Networks - NC, NS, NQ

- NC: how many classes to use per each iteration on batch (or ‘ways’)
- NS: how many examples to use as support for centroids calculus (or ‘shots’)
- NQ: how many examples to use as queries for centroids calculus (‘query’)
 - 1 shots -> One-Shot Learning
 - 5 shots -> Few-Shot Learning





Loss, optimizer, scheduler

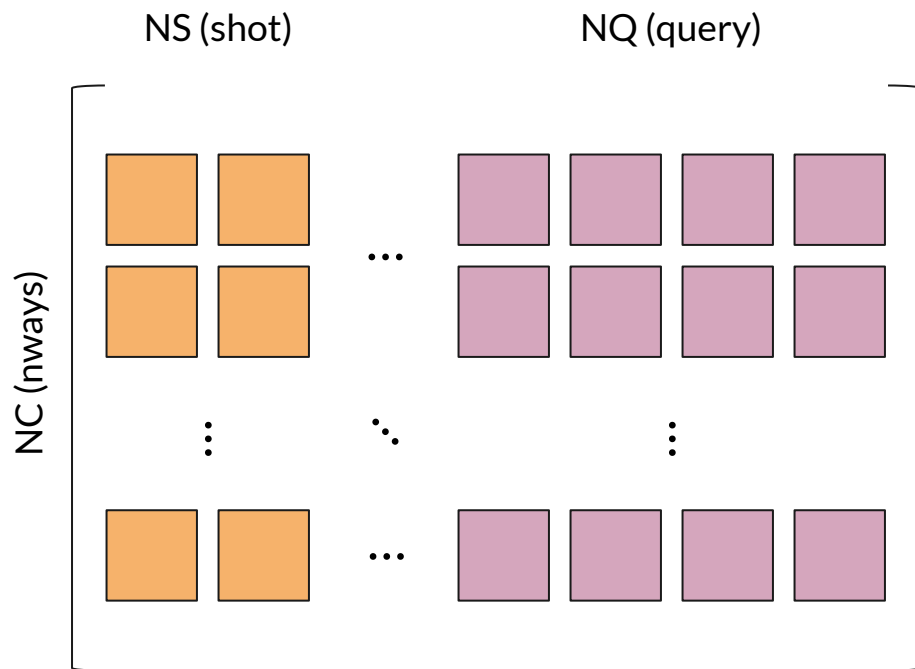
- Loss: `mean(log_softmax(-distances(query, centroids), targets)`
- Optimizer: Adam with `lr=0.001` (faster and similar results to SGD)
- Scheduler: StepLr with `step_size=20, gamma = 0.5`

Training skeleton

```
model = PrototypicalNet()
loss = loss_function(x, y, NC, NS, NQ)
optim = StepLr()
```

```
for epoch in epochs:
    for it in iterations:
        x, y = GetSample(NC, NS, NQ)
        out = model(x)
        ... loss
        backward()
    for it in iterations:
        .... eval ....
```

One batch





Code

All code is available at [github](#)

- Full hyperparams control
- 3 available datasets
- results and plots
- simple train.py and test.py scripts

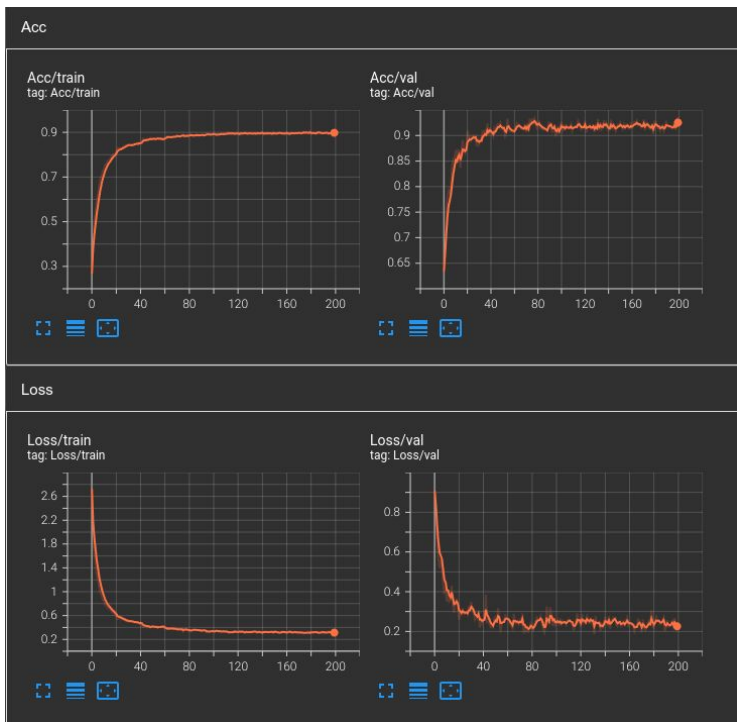
```
python train.py --dataset mini_imagenet \  
    --epochs 200 \  
    --gpu \  
    --train-num-class 30 \  
    --test-num-class 5 \  
    --number-support 5 \  
    --train-num-query 15 \  
    --episodes-per-epoch 100 \  
    --adam-lr 0.001 \  
    --opt-step-size 20 \  
    --opt-gamma 0.5 \  
    --distance-function "euclidean" \  
    --save-each 5
```



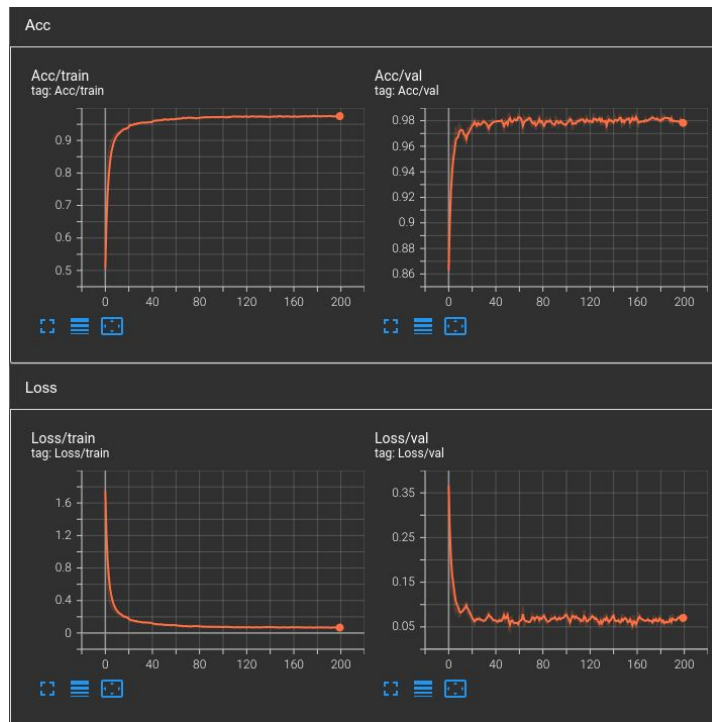
Experiments

Omniglot

1-shot

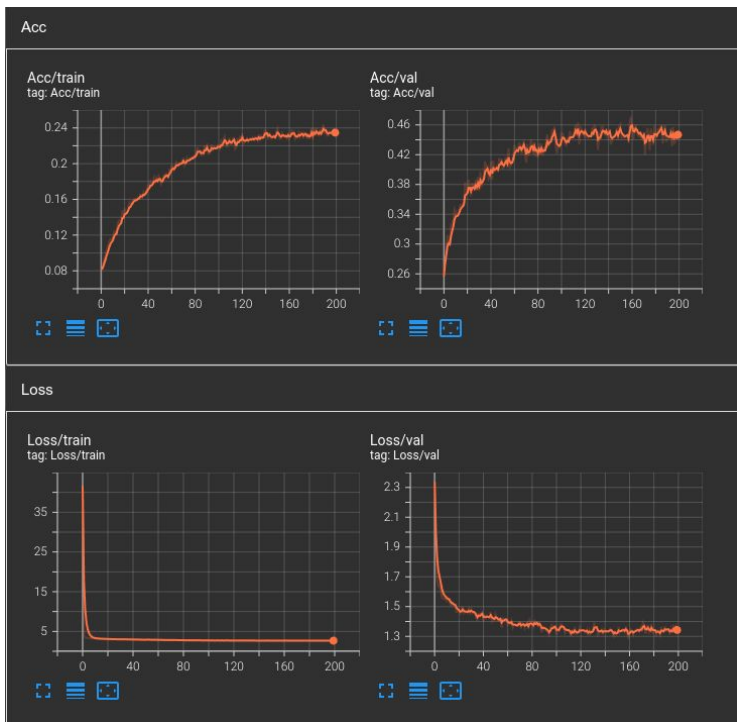


5-shot

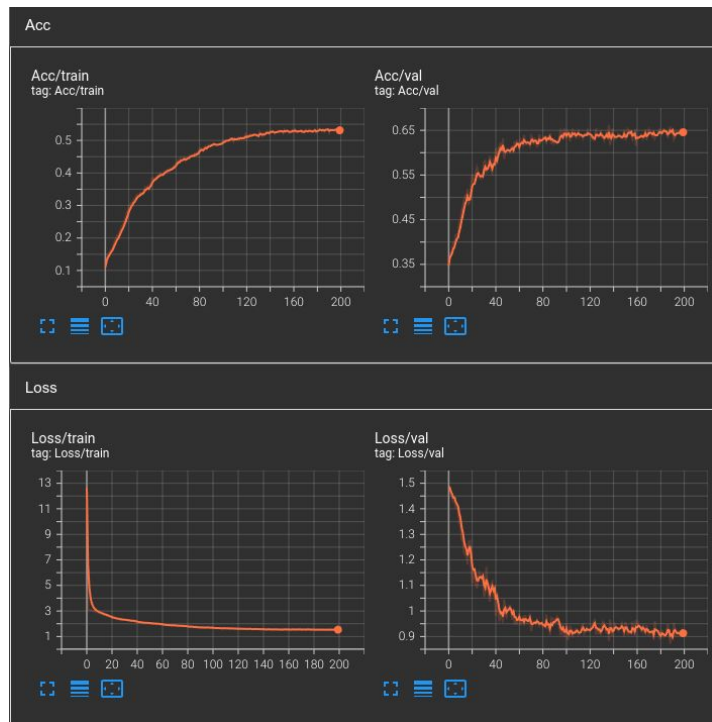


Mini Imagenet

1-shot

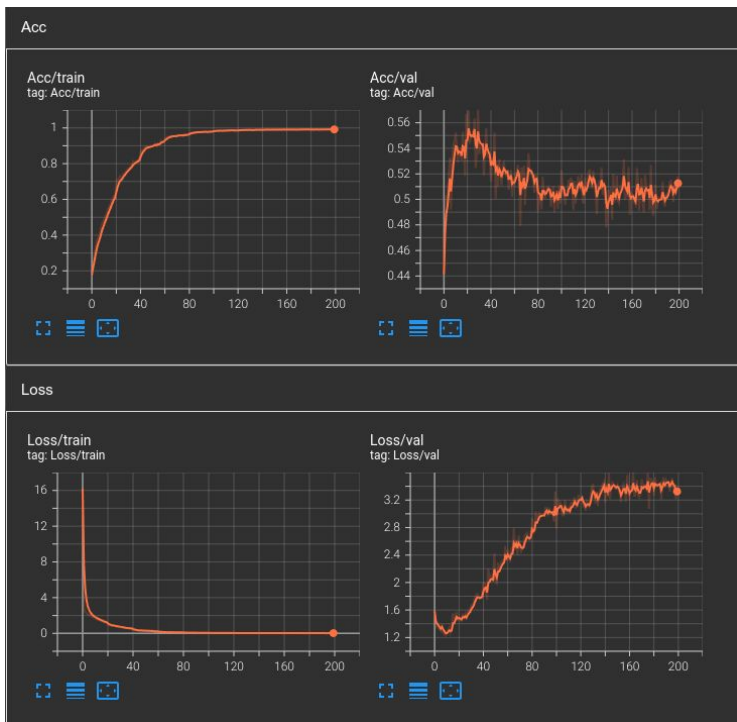


5-shot

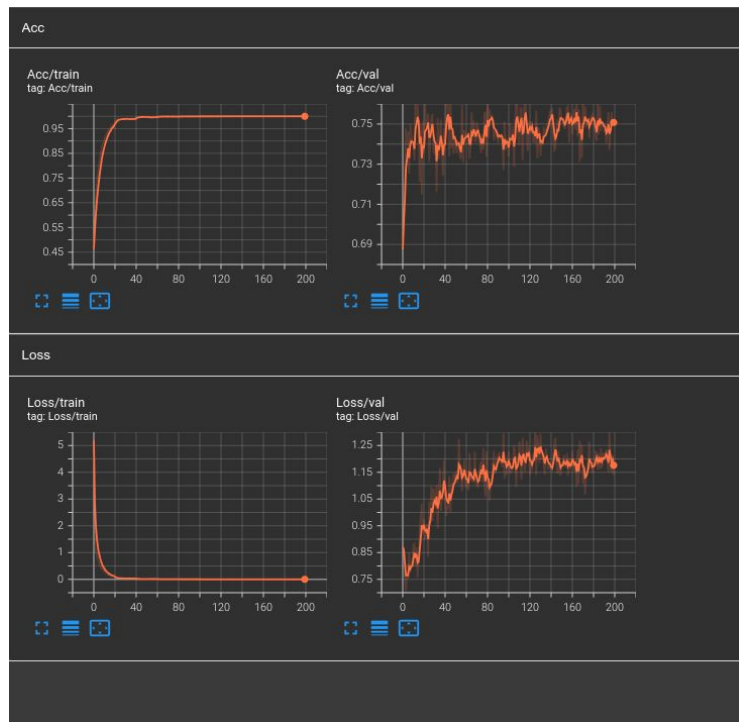


Flowers 102


1-shot



5-shot




Comparison - distance metric




Dataset	Cosine (acc)	Euclidean (acc)
mini_imagenet	22.36	63.62
omniglot	23.48	97.77
flowers102	82.89	84.48

Comparison: 5-shot



Dataset	Paper res 5-way 5-shot (Acc)	Our res 5-way 5-shot (Acc)	Paper res 5-way 1-shot (Acc)	Our res 5-way 1-shot (Acc)
mini_imagenet	68.20	63.62	49.42	46.13
omniglot	98.80	97.77	98.8	91.93
flowers102	/	84.48	/	56.08

Comparison: 1-shot



Dataset	Paper res 5-way 5-shot (Acc)	Our res 5-way 5-shot (Acc)	Paper res 5-way 1-shot (Acc)	Our res 5-way 1-shot (Acc)
mini_imagenet	68.20	63.62	49.42	46.13
omniglot	98.80	97.77	98.8	91.93
flowers102	/	84.48	/	56.08



Conclusions



Conclusion

- Euclidean distance performs better than cosine similarity
- Paper results were correctly replicated

Future studies

- Add custom dataset option for training
- Implement proper `torch.nn.Dataset` and `torch.nn.Sampler` + `torch.nn.DataLoader`
- Try different fields than CV